

Numerical Optimization

Chris Cornwell

Sept. 25, 2025

General approach to numerical optimization

Approach to minimization of function g .

General approach to numerical optimization

Approach to minimization of function g .

1. Start the process from some initial point $\mathbf{w}^{(0)}$.

General approach to numerical optimization


Approach to minimization of function g .

1. Start the process from some initial point $\mathbf{w}^{(0)}$.
2. After t steps, with $\mathbf{w}^{(t)}$ as current input to g , update it to some $\mathbf{w}^{(t)}$, “going downhill” toward a stationary point.

General approach to numerical optimization

Approach to minimization of function g .

1. Start the process from some initial point $\mathbf{w}^{(0)}$.
2. After t steps, with $\mathbf{w}^{(t)}$ as current input to g , update it to some $\mathbf{w}^{(t+1)}$, “going downhill” toward a stationary point.
3. Repeat Step (2) – converging to a stationary point, in good scenario – until you meet a **stopping condition** at some step T . The approximate stationary point (potentially a minimizer) is $\mathbf{w}^{(T)}$.



Outline

Gradient Descent

Newton's method - a second-order method

Examples and some convergence guarantees

Gradient Descent - a first-order method

Looking to minimize a loss (cost) function $\ell : \mathbb{R}^N \rightarrow \mathbb{R}$. Want to approximate stationary point for ℓ in \mathbb{R}^N , in order to minimize ℓ (hopefully).

¹For us, we will usually use α_{k+1} that is constant in t , meaning it is the same for all

Gradient Descent - a first-order method

Looking to minimize a loss (cost) function $\ell : \mathbb{R}^N \rightarrow \mathbb{R}$. Want to approximate stationary point for ℓ in \mathbb{R}^N , in order to minimize ℓ (hopefully).

- In addition to choosing initial $\mathbf{w}^{(0)} \in \mathbb{R}^N$, pick a **learning rate** $\alpha_{t+1} > 0$ for each step in the process (also called “step size”).¹

¹For us, we will usually use α_{t+1} that is constant in t , meaning it is the same for all

Gradient Descent - a first-order method

Looking to minimize a loss (cost) function $\ell : \mathbb{R}^N \rightarrow \mathbb{R}$. Want to approximate stationary point for ℓ in \mathbb{R}^N , in order to minimize ℓ (hopefully).

- ▶ In addition to choosing initial $\mathbf{w}^{(0)} \in \mathbb{R}^N$, pick a **learning rate** $\alpha_{t+1} > 0$ for each step in the process (also called “step size”).¹
- ▶ For each $t \geq 0$, iteratively assign

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \alpha_{t+1} \nabla \ell(\mathbf{w}^{(t)}).$$

¹For us, we will usually use α_{t+1} that is constant in t , meaning it is the same for all

Gradient Descent - a first-order method

Looking to minimize a loss (cost) function $\ell : \mathbb{R}^N \rightarrow \mathbb{R}$. Want to approximate stationary point for ℓ in \mathbb{R}^N , in order to minimize ℓ (hopefully).

- ▶ In addition to choosing initial $\mathbf{w}^{(0)} \in \mathbb{R}^N$, pick a **learning rate** $\alpha_{t+1} > 0$ for each step in the process (also called “step size”).¹
- ▶ For each $t \geq 0$, iteratively assign

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \alpha_{t+1} \nabla \ell(\mathbf{w}^{(t)}).$$

- ▶ As in the general optimization description, proceed until a **stopping condition** is met.

¹For us, we will usually use α_{t+1} that is constant in t , meaning it is the same for all

Gradient Descent - a first-order method

Looking to minimize a loss (cost) function $\ell : \mathbb{R}^N \rightarrow \mathbb{R}$. Want to approximate stationary point for ℓ in \mathbb{R}^N , in order to minimize ℓ (hopefully).

- ▶ In addition to choosing initial $\mathbf{w}^{(0)} \in \mathbb{R}^N$, pick a **learning rate** $\alpha_{t+1} > 0$ for each step in the process (also called “step size”).¹
- ▶ For each $t \geq 0$, iteratively assign

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \alpha_{t+1} \nabla \ell(\mathbf{w}^{(t)}).$$

- ▶ As in the general optimization description, proceed until a **stopping condition** is met.

More on stopping conditions later. For now, ...

¹For us, we will usually use α_{t+1} that is constant in t , meaning it is the same for all

Gradient Descent - a first-order method

Looking to minimize a loss (cost) function $\ell : \mathbb{R}^N \rightarrow \mathbb{R}$. Want to approximate stationary point for ℓ in \mathbb{R}^N , in order to minimize ℓ (hopefully).

- ▶ In addition to choosing initial $\mathbf{w}^{(0)} \in \mathbb{R}^N$, pick a **learning rate** $\alpha_{t+1} > 0$ for each step in the process (also called “step size”).¹
- ▶ For each $t \geq 0$, iteratively assign

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \alpha_{t+1} \nabla \ell(\mathbf{w}^{(t)}).$$

- ▶ As in the general optimization description, proceed until a **stopping condition** is met.

More on stopping conditions later. For now, ...

- ▶ choose a (“small”) threshold value ε . Stop when, as part of the last update, the change in every parameter divided by its size is not more than ε .

¹For us, we will usually use α_{t+1} that is constant in t , meaning it is the same for all

Gradient Descent - a first-order method

Looking to minimize a loss (cost) function $\ell : \mathbb{R}^N \rightarrow \mathbb{R}$. Want to approximate stationary point for ℓ in \mathbb{R}^N , in order to minimize ℓ (hopefully).

- ▶ In addition to choosing initial $\mathbf{w}^{(0)} \in \mathbb{R}^N$, pick a **learning rate** $\alpha_{t+1} > 0$ for each step in the process (also called “step size”).¹
- ▶ For each $t \geq 0$, iteratively assign

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \alpha_{t+1} \nabla \ell(\mathbf{w}^{(t)}).$$

- ▶ As in the general optimization description, proceed until a **stopping condition** is met.

More on stopping conditions later. For now, ...

- ▶ choose a (“small”) threshold value ε . Stop when, as part of the last update, the change in every parameter divided by its size is not more than ε . That is, stop when

$$\frac{|\omega_j^{(t)} - \omega_j^{(t-1)}|}{|\omega_j^{(t-1)}|} \leq \varepsilon, \quad \forall 1 \leq j \leq p.$$

¹For us, we will usually use α_{t+1} that is constant in t , meaning it is the same for all

Relation to Linear approximations

The update step in Gradient Descent uses the linear approximation to the loss function in order to make an update.

Relation to Linear approximations

The update step in Gradient Descent uses the linear approximation to the loss function in order to make an update.

Note: linear approximation is

$$h(\mathbf{w}) = \ell(\mathbf{w}^{(t)}) + \nabla \ell(\mathbf{w}^{(t)})^T (\mathbf{w} - \mathbf{w}^{(t)}).$$

How does one change $\mathbf{w}^{(t)}$ to make $h(\mathbf{w})$ decrease the most? Want some (fixed length vector) $\Delta \mathbf{w}$ so that $h(\mathbf{w}^{(t)}) - h(\mathbf{w}^{(t)} + \Delta \mathbf{w})$ is as large as possible.

Relation to Linear approximations

The update step in Gradient Descent uses the linear approximation to the loss function in order to make an update.

Note: linear approximation is

$$h(\mathbf{w}) = \ell(\mathbf{w}^{(t)}) + \nabla \ell(\mathbf{w}^{(t)})^T (\mathbf{w} - \mathbf{w}^{(t)}).$$

How does one change $\mathbf{w}^{(t)}$ to make $h(\mathbf{w})$ decrease the most? Want some (fixed length vector) $\Delta \mathbf{w}$ so that $h(\mathbf{w}^{(t)}) - h(\mathbf{w}^{(t)} + \Delta \mathbf{w})$ is as large as possible.

Since

$$h(\mathbf{w}^{(t)}) - h(\mathbf{w}^{(t)} + \Delta \mathbf{w}) = -\nabla \ell(\mathbf{w}^{(t)})^T \Delta \mathbf{w}$$

this occurs when $\nabla \ell(\mathbf{w}^{(t)})^T \Delta \mathbf{w}$ is as negative as possible, which is when $\Delta \mathbf{w}$ is in the opposite direction of $\nabla \ell(\mathbf{w}^{(t)})$; and so, our update is $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \Delta \mathbf{w} = \mathbf{w}^{(t)} - \alpha_{t+1} \nabla \ell(\mathbf{w}^{(t)})$.

Convergence

Does gradient descent converge to point at which loss function is minimized? Is it even guaranteed to converge?

Convergence

Does gradient descent converge to point at which loss function is minimized? Is it even guaranteed to converge?

Short answer: No, not necessarily.

...so, *in what cases* can we guarantee such a thing?

Issues with convergence

To demonstrate the difficulty, imagine a “toy” loss function: $\ell : \mathbb{R} \rightarrow \mathbb{R}$ with $\ell(w) = w^2$. At each w we have $\nabla \ell = \left(\frac{d\ell}{dw} \right) = (2w)$.

Issues with convergence

To demonstrate the difficulty, imagine a “toy” loss function: $\ell : \mathbb{R} \rightarrow \mathbb{R}$ with $\ell(w) = w^2$. At each w we have $\nabla \ell = \left(\frac{d\ell}{dw} \right) = (2w)$.

Say learning rate: $\alpha_{t+1} > 1$ for all t . Then, at any $w^{(t)} > 0$, we get

$$w^{(t+1)} = w^{(t)} - 2\alpha_{t+1}w^{(t)} < w^{(t)} - 2w^{(t)} = -w^{(t)},$$

and so $w^{(t+1)} < -w^{(t)}$ which means that $|w^{(t+1)}| > |w^{(t)}|$.

Issues with convergence

To demonstrate the difficulty, imagine a “toy” loss function: $\ell : \mathbb{R} \rightarrow \mathbb{R}$ with $\ell(w) = w^2$. At each w we have $\nabla \ell = \left(\frac{d\ell}{dw} \right) = (2w)$.

Say learning rate: $\alpha_{t+1} > 1$ for all t . Then, at any $w^{(t)} > 0$, we get

$$w^{(t+1)} = w^{(t)} - 2\alpha_{t+1}w^{(t)} < w^{(t)} - 2w^{(t)} = -w^{(t)},$$

and so $w^{(t+1)} < -w^{(t)}$ which means that $|w^{(t+1)}| > |w^{(t)}|$. Likewise, if $w^{(t)} < 0$ then $|w^{(t+1)}| > |w^{(t)}|$.

Issues with convergence

To demonstrate the difficulty, imagine a “toy” loss function: $\ell : \mathbb{R} \rightarrow \mathbb{R}$ with $\ell(w) = w^2$. At each w we have $\nabla \ell = \left(\frac{d\ell}{dw} \right) = (2w)$.

Say learning rate: $\alpha_{t+1} > 1$ for all t . Then, at any $w^{(t)} > 0$, we get

$$w^{(t+1)} = w^{(t)} - 2\alpha_{t+1}w^{(t)} < w^{(t)} - 2w^{(t)} = -w^{(t)},$$

and so $w^{(t+1)} < -w^{(t)}$ which means that $|w^{(t+1)}| > |w^{(t)}|$. Likewise, if $w^{(t)} < 0$ then $|w^{(t+1)}| > |w^{(t)}|$.

So, get divergence when $\alpha_{t+1} > 1$. However, if $0 < \alpha_{t+1} < 1$, then (for the function $\ell(w) = w^2$, at least) it will converge to minimizer $w = 0$.

Issues with convergence

To demonstrate the difficulty, imagine a “toy” loss function: $\ell : \mathbb{R} \rightarrow \mathbb{R}$ with $\ell(w) = w^2$. At each w we have $\nabla \ell = \left(\frac{d\ell}{dw} \right) = (2w)$.

Say learning rate: $\alpha_{t+1} > 1$ for all t . Then, at any $w^{(t)} > 0$, we get

$$w^{(t+1)} = w^{(t)} - 2\alpha_{t+1}w^{(t)} < w^{(t)} - 2w^{(t)} = -w^{(t)},$$

and so $w^{(t+1)} < -w^{(t)}$ which means that $|w^{(t+1)}| > |w^{(t)}|$. Likewise, if $w^{(t)} < 0$ then $|w^{(t+1)}| > |w^{(t)}|$.

So, get divergence when $\alpha_{t+1} > 1$. However, if $0 < \alpha_{t+1} < 1$, then (for the function $\ell(w) = w^2$, at least) it will converge to minimizer $w = 0$.




Figure: Gradient descent on $\ell(w) = w^2$. Left: $\alpha = 1.05$; right: $\alpha = 0.95$.

Outline

Gradient Descent

Newton's method - a second-order method

Examples and some convergence guarantees

Newton's method

As before, we want to use approximations to get closer to a stationary point of $\ell : \mathbb{R}^N \rightarrow \mathbb{R}$, however we will now use a second order approximation:

$$h(\mathbf{w}) = \ell(\mathbf{w}^{(t)}) + \nabla \ell(\mathbf{w}^{(t)})^T (\mathbf{w} - \mathbf{w}^{(t)}) + \frac{1}{2} (\mathbf{w} - \mathbf{w}^{(t)})^T \nabla^2 \ell(\mathbf{w}^{(t)}) (\mathbf{w} - \mathbf{w}^{(t)}).$$

Newton's method

As before, we want to use approximations to get closer to a stationary point of $\ell : \mathbb{R}^N \rightarrow \mathbb{R}$, however we will now use a second order approximation:

$$h(\mathbf{w}) = \ell(\mathbf{w}^{(t)}) + \nabla \ell(\mathbf{w}^{(t)})^T (\mathbf{w} - \mathbf{w}^{(t)}) + \frac{1}{2} (\mathbf{w} - \mathbf{w}^{(t)})^T \nabla^2 \ell(\mathbf{w}^{(t)}) (\mathbf{w} - \mathbf{w}^{(t)}).$$

As a consequence from the recent homework, this has a stationary point at a solution to the linear equation

$\nabla^2 \ell(\mathbf{w}^{(t)}) \mathbf{w} = \nabla^2 \ell(\mathbf{w}^{(t)}) \mathbf{w}^{(t)} - \nabla \ell(\mathbf{w}^{(t)})$. Call that solution $\mathbf{w}^{(t+1)}$.

Repeat the above step until some stopping condition is met.

Newton's method - advantages and disadvantages

Advantage: When converging to a minimum of the function, Newton's method will converge in fewer steps.

Disadvantages: One has to compute the Hessian – more computation time and more storage needed. Solving the linear system can also have numerical issues and take more time.

Also, if ℓ is not convex (and $\mathbf{w}^{(t)}$ is not in a convex region for the function), then it may approach a maximum instead of minimum.

Outline

Gradient Descent

Newton's method - a second-order method

Examples and some convergence guarantees

Example with Linear regression

Last time: given sample data \mathcal{S} for simple linear regression, and using MSE as empirical loss, $\mathcal{L}_{\mathcal{S}}(m, b) = \frac{1}{n} \sum_{i=1}^n (mx_i + b - y_i)^2$, we found

$$\nabla \mathcal{L}_{\mathcal{S}}(m, b) = \left(\frac{2}{n} \sum_{i=1}^n (mx_i + b - y_i)x_i, \quad \frac{2}{n} \sum_{i=1}^n (mx_i + b - y_i) \right).$$

Example with Linear regression


Last time: given sample data \mathcal{S} for simple linear regression, and using MSE as empirical loss, $\mathcal{L}_{\mathcal{S}}(m, b) = \frac{1}{n} \sum_{i=1}^n (mx_i + b - y_i)^2$, we found

$$\nabla \mathcal{L}_{\mathcal{S}}(m, b) = \left(\frac{2}{n} \sum_{i=1}^n (mx_i + b - y_i)x_i, \quad \frac{2}{n} \sum_{i=1}^n (mx_i + b - y_i) \right).$$

Example: batch gradient descent working on the '[Example1.csv](#)' data.

The LSR line, using closed form.

► $\hat{m} \approx 1.520$, $\hat{b} = -0.3346$:




Example with Linear regression

Last time: given sample data \mathcal{S} for simple linear regression, and using MSE as empirical loss, $\mathcal{L}_{\mathcal{S}}(m, b) = \frac{1}{n} \sum_{i=1}^n (mx_i + b - y_i)^2$, we found

$$\nabla \mathcal{L}_{\mathcal{S}}(m, b) = \left(\frac{2}{n} \sum_{i=1}^n (mx_i + b - y_i)x_i, \frac{2}{n} \sum_{i=1}^n (mx_i + b - y_i) \right).$$

Example: batch gradient descent working on the 'Example1.csv' data.

Plot of selected lines found during batch GD updates; starting parameters $m = 0, b = 0$;




Example with Linear regression

Last time: given sample data \mathcal{S} for simple linear regression, and using MSE as empirical loss, $\mathcal{L}_{\mathcal{S}}(m, b) = \frac{1}{n} \sum_{i=1}^n (mx_i + b - y_i)^2$, we found

$$\nabla \mathcal{L}_{\mathcal{S}}(m, b) = \left(\frac{2}{n} \sum_{i=1}^n (mx_i + b - y_i)x_i, \frac{2}{n} \sum_{i=1}^n (mx_i + b - y_i) \right).$$

Example: batch gradient descent working on the 'Example1.csv' data.

Plot of selected lines found during batch GD updates; starting parameters $m = 0, b = 0$; learning rate set to 0.1.
Parameter values on iteration 208: $m = 1.519, b = -0.334$.



Implementing Gradient Descent

Some notes on implementation of the gradient descent updates.

1. Each partial derivative of $\mathcal{L}_{\mathcal{S}}$, can compute it in one line of code.

Implementing Gradient Descent

Some notes on implementation of the gradient descent updates.

1. Each partial derivative of \mathcal{L}_S , can compute it in one line of code.

Presuming x and y are arrays of coordinates for sample data, and that

$n = \text{len}(x)$, the following computes the partial derivatives:

```
1 | partial_m = (2/n)*np.sum( (m*x + b - y)*x )  
2 | partial_b = (2/n)*np.sum( (m*x + b - y) )
```

Implementing Gradient Descent

Some notes on implementation of the gradient descent updates.

1. Each partial derivative of \mathcal{L}_S , can compute it in one line of code.

Presuming x and y are arrays of coordinates for sample data, and that $n = \text{len}(x)$, the following computes the partial derivatives:

```
1 | partial_m = (2/n)*np.sum( (m*x + b - y)*x )  
2 | partial_b = (2/n)*np.sum( (m*x + b - y) )
```

2. To implement GD, want more than one function – at the least, one to compute the gradient (given current parameters); another that performs update and checks for stopping. *Roughly...*

Implementing Gradient Descent

Some notes on implementation of the gradient descent updates.

1. Each partial derivative of \mathcal{L}_S , can compute it in one line of code.

Presuming x and y are arrays of coordinates for sample data, and that $n = \text{len}(x)$, the following computes the partial derivatives:

```
1 | partial_m = (2/n)*np.sum( (m*x + b - y)*x )
2 | partial_b = (2/n)*np.sum( (m*x + b - y) )
```

2. To implement GD, want more than one function – at the least, one to compute the gradient (given current parameters); another that performs update and checks for stopping. *Roughly...*

```
## lr is learning rate; threshold is for stopping;
input: X, y, lr, threshold
p ← initial array of parameters
while (max of last_update > threshold){
    grad ← compute_grad(p, X, y)
    last_update ← | grad / p | ## entrywise array division
    # handle p[i] near 0
    p ← p - lr*grad
}
return p
```

Guarantees of convergence

If your loss function is differentiable and a **convex function**, and if have some “control” on size of the gradient then, by choosing η small enough, can guarantee convergence.

²Meaning: \exists a constant C s.t. for all ω_1, ω_2 , $|\nabla \mathcal{L}(\omega_1) - \nabla \mathcal{L}(\omega_2)| \leq C|\omega_1 - \omega_2|$.

Guarantees of convergence

If your loss function is differentiable and a **convex function**, and if you have some “control” on size of the gradient then, by choosing η small enough, you can guarantee convergence.

Theorem

Suppose that $\mathcal{L} : \mathbb{R}^p \rightarrow \mathbb{R}$ is differentiable and convex, and suppose that $\nabla \mathcal{L}$ is Lipschitz continuous² with some constant $C > 0$ and that $\eta \leq 1/C$. Then, for a minimizer ω^* of \mathcal{L} ,

$$\mathcal{L}(\omega^{(t)}) - \mathcal{L}(\omega^*) \leq \frac{|\omega^{(0)} - \omega^*|^2}{2\eta t}.$$

²Meaning: \exists a constant C s.t. for all ω_1, ω_2 , $|\nabla \mathcal{L}(\omega_1) - \nabla \mathcal{L}(\omega_2)| \leq C|\omega_1 - \omega_2|$.

Guarantees of convergence

If your loss function is differentiable and a **convex function**, and if you have some “control” on size of the gradient then, by choosing η small enough, you can guarantee convergence.

Theorem

Suppose that $\mathcal{L} : \mathbb{R}^p \rightarrow \mathbb{R}$ is differentiable and convex, and suppose that $\nabla \mathcal{L}$ is Lipschitz continuous² with some constant $C > 0$ and that $\eta \leq 1/C$. Then, for a minimizer ω^* of \mathcal{L} ,

$$\mathcal{L}(\omega^{(t)}) - \mathcal{L}(\omega^*) \leq \frac{|\omega^{(0)} - \omega^*|^2}{2\eta t}.$$

- The difference between $\mathcal{L}(\omega^{(t)})$ and the minimum of \mathcal{L} is bounded by a constant times $1/t$.

²Meaning: \exists a constant C s.t. for all ω_1, ω_2 , $|\nabla \mathcal{L}(\omega_1) - \nabla \mathcal{L}(\omega_2)| \leq C|\omega_1 - \omega_2|$.

Guarantees of convergence

Previous theorem requires using actual gradient of \mathcal{L} in each update step. Here is a convergence guarantee that allows for a random vector \mathbf{D}_t , in place of $\nabla \mathcal{L}$, as long as $\mathbb{E}[\mathbf{D}_t | \omega^{(t)}] = \nabla \mathcal{L}(\omega^{(t)})$.

- e.g., in Stochastic or Mini-batch gradient descent.

Guarantees of convergence

Previous theorem requires using actual gradient of \mathcal{L} in each update step. Here is a convergence guarantee that allows for a random vector \mathbf{D}_t , in place of $\nabla \mathcal{L}$, as long as $\mathbb{E}[\mathbf{D}_t | \omega^{(t)}] = \nabla \mathcal{L}(\omega^{(t)})$.

► e.g., in Stochastic or Mini-batch gradient descent.

Theorem

Suppose that \mathcal{L} is differentiable and convex, that $\omega^{(0)} = \mathbf{0}$, and that $\eta = \frac{1}{\sqrt{K}}$ for an integer $K > 0$. Finally, suppose that $|\mathbf{D}_t| \leq 1$ for all $1 \leq t \leq K$. Then, for a minimizer ω^* of \mathcal{L} ,

$$\mathbb{E}[\mathcal{L}(\bar{\omega})] - \mathcal{L}(\omega^*) \leq \frac{1}{\sqrt{K}}$$

where $\bar{\omega}$ is the average of $\omega^{(1)}, \omega^{(2)}, \dots, \omega^{(K)}$.